Azure, as a cloud platform, has a lot of storage services, including various SQL databases. This document contains a comparison table with performance results for multi-model database service CosmosDB, using its SQL API, and few most popular SQL databases.

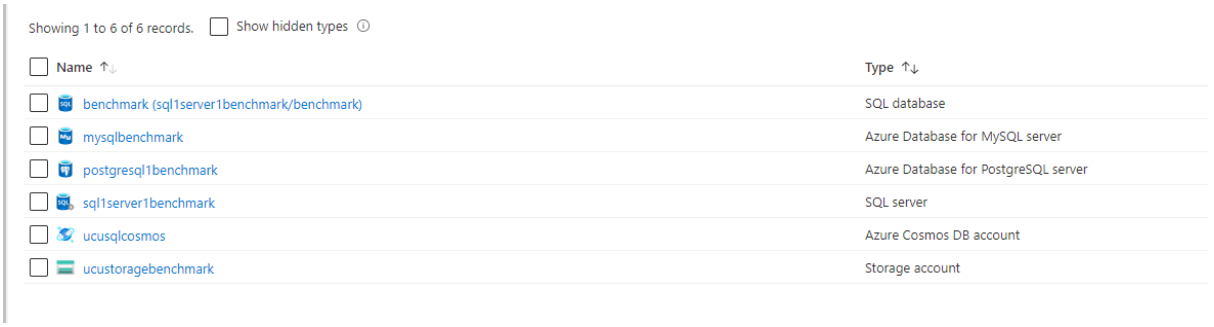

Code in the repository contains Jupiter notebooks to populate databases with test data.

Test Data	**CosmosDB (SQL API)**	**PosgreSQL**	**SQL Server**	**MySQL**
Dataset: file with 48439 rows, 21 columns, 4.4 MB https://globaldatalab.org/assets/2019/09/SHDI%20Complete%203.0.csv	Insert from Python app - near 15 minutes Read data - 7.4 sec	Insert into table - 45.1 s Read data - 59.3 s	Insert into table - 52 s Read data - 45 s	Insert into table - 55 s Read data - 70 s

CosmosDB (SQL API) being the most expensive has no reliable Python client. Because of that inserting consumes a lot of time using the loop approach. Complexity is O(n^). This DB service shows better results on high-concurrent workloads, rather than huge bulk data insertion.
Despite the loop approach for inserting data, reading was fast, comparing to other database services.
The lowest possible capacity was used for CosmosDB (Throughput (RU/s) - 400).

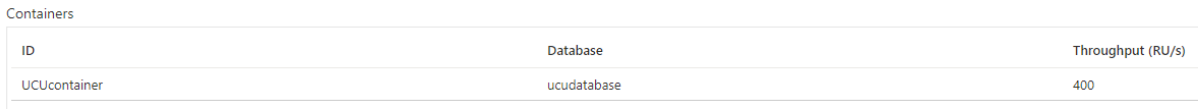

PostgreSQL, SQL Server, MySQL showed near the same result, however, PostgreSQL was slightly faster.
Those databases have comparably similar capacities at the lowest level.